

## 1 Code Explanation

The answer to exercise 1 is in the comment of `h_timer.c`.

## 2 Measure Methods and Results

In order to measure the time precisely without the overhead of the initialization of counter and for loop. I first loop for 100000 times and initialize two counters each time in for loop and perform the operation between two counters and then I just loop for 100000 times and initialize two counters each time in for loop and do nothing between the two counters. Then I subtract the two results and divide the difference by the number of iteration.

### 2.1 Function Call

The time for a minimum function call is 0.1 ns in a csil machine and 0.4 - 0.6 ns in my laptop.

I first call function `clock_gettime()` with a `CLOCK_PROCESS_CPUTIME_ID` counter to get the start time. And use a for loop to test the empty function call for 100000 times. And then call function `clock_gettime()` with a `CLOCK_PROCESS_CPUTIME_ID` counter to get the stop time. Result is computed by subtraction of the start and stop.

And in order to reduce the overhead of initialization of clock and for loop, I also call function `clock_gettime()` twice and do an idle for loop for 100000 times.

### 2.2 System Call

The time for a minimum function call is 2-3 ns in a csil machine and my laptop.

I first call function `clock_gettime()` with a `CLOCK_PROCESS_CPUTIME_ID` counter to get the start time. And use a for loop to test the `getpid()` which is the simplest system call for 100000 times. And then call function `clock_gettime()` with a `CLOCK_PROCESS_CPUTIME_ID` counter to get the stop time. Result is computed by subtraction of the start and stop.

And in order to reduce the overhead of initialization of clock and for loop, I also call function `clock_gettime()` twice and do an idle for loop for 100000 times.

### 2.3 Process Context Switch

The time for a minimum function call is 1000-2000 ns in my laptop.

I first use `sched_setaffinity()` to set the running CPU to be one single CPU. Then, I call function `clock_gettime()` with a `CLOCK_MONOTONIC` counter to get the start time. I use `sched_yield()` to complete context switch for 100000 times. And then I call function `clock_gettime()` with a `CLOCK_MONOTONIC` counter to get the stop time. Result is computed by subtraction of the start and stop.

And in order to reduce the overhead of initialization of clock and for loop, I also call function `clock_gettime()` twice and do an idle for loop for 100000 times.

### 2.4 Thread Context Switch

The time for a minimum function call is 1000-2000 ns in my laptop.

I first use `sched_setaffinity()` to set the running CPU to be one single CPU. Then, I call function `clock_gettime()` with a `CLOCK_PROCESS_CPUTIME_ID` counter to get the start time. I use functions `pthread_mutex_lock()`, `pthread_cond_wait()` and `pthread_cond_signal()` to implement the method described in a2.pdf to complete context switch for 100000 times. And then I call function `clock_gettime()` with a `CLOCK_PROCESS_CPUTIME_ID` counter to get the stop time. Result is computed by subtraction of the start and stop.

And in order to reduce the overhead of initialization of clock and for loop, I do an for loop for 100000 times and call function *clock\_gettime()* twice in the for loop.